

Assignment3

February 7, 2023

1 Assignment 3 Group 85

Authors: *Niclas Lindmark, Noa Sjöstrand, Anton Johansson*

Course: *DAT405 Introduction to data science and AI*

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import seaborn as sns

from sklearn.cluster import DBSCAN
from sklearn import metrics
from sklearn import datasets
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import NearestNeighbors
from sklearn.metrics import silhouette_score
```

```
[39]: # reading data

d = pd.read_csv('./data_assignment3.csv')

xValues = d['phi']
yValues = d['psi']

d_PRO = pd.DataFrame(d.loc[d['residue name'] == 'PRO'])
d
```

```
[39]:
```

	residue name	position	chain	phi	psi
0	LYS	10	A	-149.312855	142.657714
1	PRO	11	A	-44.283210	136.002076
2	LYS	12	A	-119.972621	-168.705263
3	LEU	13	A	-135.317212	137.143523
4	LEU	14	A	-104.851467	95.928520
...
29364	GLY	374	B	-147.749557	155.223562
29365	GLN	375	B	-117.428541	133.019506

29366	ILE	376	B	-113.586448	112.091970
29367	ASN	377	B	-100.668779	-12.102821
29368	LYS	378	B	-169.951240	94.233680

[29369 rows x 5 columns]

1.1 1. Visualizing distribution


Show the distribution of phi and psi combinations using a scatter plot and 2D histogram

1.1.1 Scatter plot

```
[89]: plt.style.use('seaborn')
plt.rcParams["figure.figsize"] = [10, 5]
plt.rcParams["figure.autolayout"] = True

plt.scatter(xValues,yValues,s=0.5, alpha=0.8, c='orange')
plt.xlabel('phi')
plt.ylabel('psi')

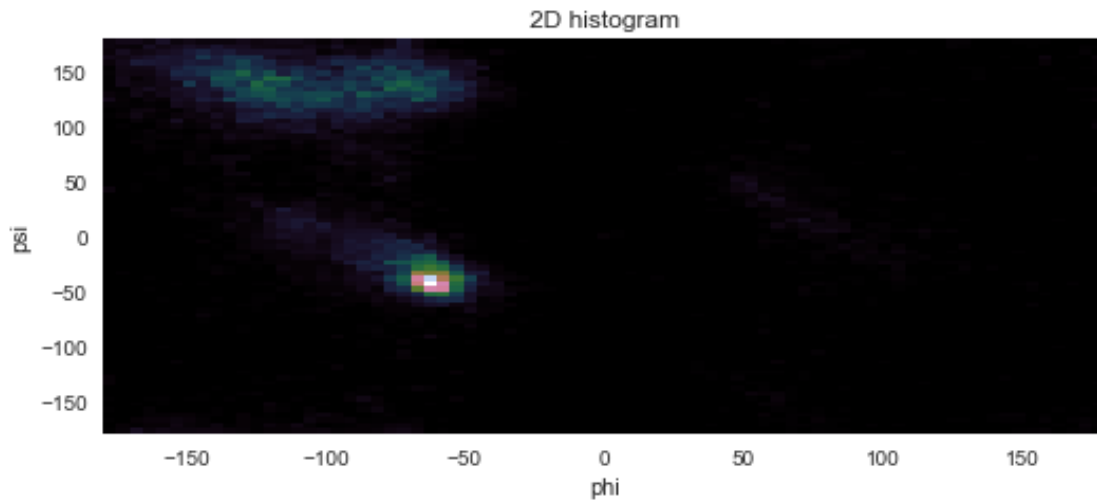
plt.show()
```



1.1.2 2D histogram

```
[73]: plt.hist2d(xValues, yValues, 75, cmap="cubehelix")
plt.xlabel('phi')
plt.ylabel('psi')
```

```
plt.title('2D histogram')
plt.show()
```



1.2 2. KMeans

Use the K-means clustering method to cluster the phi and psi angle combinations in the data file.


```
[88]: xy_vals = list(zip(xValues, yValues))
      X = np.array(xy_vals)

      score = -1
      best_fit = (0, 0, 0)
      for n_clusters in range(2, 5):
          kmeans = KMeans(n_clusters, n_init=10).fit_predict(X)
          new_score = silhouette_score(X, kmeans)
          if new_score > score:
              score = new_score
              best_fit = (n_clusters, score)
              d['color_kmeans'] = kmeans

      plt.scatter(d['phi'], d['psi'], c=d['color_kmeans'], s=0.7, cmap='jet')
      plt.xlabel('phi')
      plt.ylabel('psi')
      plt.title('KMeans')

      print(f'best fit is: {best_fit}')
      plt.show()
```

best fit is: (3, 0.6724895253169637)




```
[77]: wcss = []

for i in range(1, 11):
    clustering = KMeans(n_clusters=i, init='k-means++', random_state=42,
    ↪n_init=10)
    clustering.fit(X)
    wcss.append(clustering.inertia_)

ks = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
sns.lineplot(x = ks, y = wcss)
plt.title('Elbow')
plt.xlabel('k - value')
plt.ylabel('inertia')

plt.show()
```



1.2.1 Answer 2:


- a) We chose $k=3$ based on the elbow diagram and silhouette score. First we used the elbow diagram to approximate the k value, then we used the silhouette score to iterate through values close to our approximation to find the k value with the highest score.
- b) Our clusters seem fairly reasonable since they have separated the most obvious clusters. However, as it is a k -means clustering, all the noise has also been assigned to a cluster. That is not ideal. Also, the angles at the bottom and top of the scatter plot are actually next to each other, but our k -means clustering has not managed to take this into account.

1.3 3. DBSCAN

Using the DBSCAN method to cluster the phi and psi angle combinations in the data file.

```
[96]: neigh = NearestNeighbors(n_neighbors=2)
      nbrs = neigh.fit(X)
      distances, indices = nbrs.kneighbors(X)
      distances = np.sort(distances, axis=0)
      distances = distances[:,1]
      plt.plot(distances)
```

```
[96]: [<matplotlib.lines.Line2D at 0x7f7f4608efa0>]
```



[86]: *# computing the best DBSCAN*

```
best_e = 0
best_l = 0
best_sc = 0
best_ms = 0

for e in range(1, 50, 1):
    for ms in range(10, 300, 10):
        labels = DBSCAN(eps=e, min_samples=ms).fit_predict(X)
        if len(set(labels)) > 1:
            sc = silhouette_score(X, labels)
            if sc > best_sc:
                best_sc = sc
                best_e = e
                best_l = labels
                best_ms = ms

d['color_dbscan'] = best_l
print(f'best epsilon: {best_e}')
print(f'best min samples: {best_ms}')
print(best_sc)
print(f'Number of clusters: {len(set(best_l))-1}')


plt.scatter(d['phi'], d['psi'], c=d['color_dbscan'], s=0.7, cmap='jet')
plt.title(f'Estimated number of clusters: {len(set(best_l))-1}')
plt.xlabel('phi')
plt.ylabel('psi')
```

```

best epsilon: 15
best min samples: 200
0.6084435508126602
Number of clusters: 3

```

```
[86]: Text(0, 0.5, 'psi')
```



1.3.1 Assessing clusters and outliers

```
[90]: # counting the number of clusters
n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)
n_noise_ = list(labels).count(-1)

print('Estimated number of clusters: %d' % n_clusters_)
print('Estimated number of noise points: %d' % n_noise_)
```


```

Estimated number of clusters: 3
Estimated number of noise points: 2698

```


1.3.2 Outliers

```
[92]: d_outliers = pd.DataFrame(d.loc[d['color_dbscan'] == -1])
d_outliers['residue name'].value_counts().sort_index().plot(kind='bar', rot=0,
    ↳ ylabel='nbr of outliers', )
plt.show()
```



```
[93]: d_GLY = pd.DataFrame(d.loc[d['residue name'] == 'GLY'])
plt.scatter(d_GLY['phi'], d_GLY['psi'], c=d_GLY['color_dbscan'], s=0.
→7,cmap='jet')
```

```
[93]: <matplotlib.collections.PathCollection at 0x7f7f39b1b700>
```



1.3.3 Answer 3:

- a) We chose eps and minimum num of sample points based on the silhouette score. First we used nearest_neighbors which gave 3 as the optimal eps. But after trying eps=3 at different

min_samples, no matter the value on minimum num of samples the silhouette score was poor (under 0.3).

We iterated through different values for epsilon and min num samples to understand what was reasonable. We did not want more than 5 clusters, not too much noise and a high silhouette score. We iterated through different values for eps and min samples and found that eps could be 15 and min sample should be 200. That combination achieved a silhouette score of 0.608, which is acceptable, whilst keeping the number of clusters at 3.

- b) You can see the clusters in our plots.
- c) You can see this in the graphs above.


1.4 4. DBSCAN with only residue type PRO

The data file can be stratified by amino acid residue type. Use DBSCAN to cluster the data that have residue type PRO.

1.4.1 Plot data

```
[99]: d_PRO = pd.DataFrame(d.loc[d['residue name'] == 'PRO'])  
plt.scatter(d_PRO['phi'], d_PRO['psi'], s=1, c='brown')
```

```
[99]: <matplotlib.collections.PathCollection at 0x7f801bf520>
```



1.4.2 DBSCAN on PRO data

```
[100]: xy_PROvals = list(zip(d_PRO['phi'], d_PRO['psi']))  
XPRO = np.array(xy_PROvals)  
  
# computing the DBSCAN
```

```

best_e = 0
best_l = 0
best_sc = 0
best_ms = 0
for e in range(30, 55, 1):
    for ms in range(50, 500, 10):
        labels = DBSCAN(eps=e, min_samples=ms).fit_predict(XPRO)
        if len(set(labels)) > 1:
            sc = silhouette_score(XPRO, labels)
            if sc > best_sc:
                best_sc = sc
                best_e = e
                best_l = labels
                best_ms = ms

d_PRO['color_dbscan'] = best_l
print(f'best epsilon: {best_e}')
print(f'best min samples: {best_ms}')
print(best_sc)
d_PRO['color_dbscan'] = best_l
n_clusters_ = len(set(best_l)) - (1 if -1 in best_l else 0)
n_noise_ = list(best_l).count(-1)
print('Estimated number of clusters: %d' % n_clusters_)
print('Estimated number of noise points: %d' % n_noise_)

plt.scatter(d_PRO['phi'], d_PRO['psi'], c=d_PRO['color_dbscan'], s=0.
    →7, cmap='jet')
plt.title('Estimated number of clusters: %d' % n_clusters_)
print(len(set(best_l)))

```

```

best epsilon: 50
best min samples: 250
0.7782783299593291
Estimated number of clusters: 2
Estimated number of noise points: 11
3

```



1.4.3 Answer 4:

Compared to the general clusters the PRO variants are more easily “clusterable”, meaning there is less noise. We have two distinct areas of clusters with high density of this variant. Noise points went from over 2500 to 11. Even though the sample size is smaller in the only-PRO data this is still a considerable difference between outliers in relation to the number of samples.

This also resulted in completely different optimal values for epsilon and min num of samples. Both values were bigger. Also, a better silhouette score indicates the clustering is more accurate.